

Assignment 9

Reid Ginoza

4/12/2020

Analytical Maximum Likelihood

This is Exercise 11.1 from (Särkkä and Solin 2019). We are given the following:

$$x_k = \theta + r_k, \quad k = 1, 2, \dots, T, \quad (1)$$

where $r_k \sim N(0, \sigma^2)$ are independent.

Maximum Likelihood Estimator

Let $L(\cdot)$ be the likelihood and $l(\cdot)$ be the negative log likelihood. Then, to find the maximum likelihood estimator $\theta_{\text{ML}} = \arg \min_{\theta} L(\theta)$, we will minimize $l(\theta)$:

$$x_k = \theta + r_k \sim N(\theta, \sigma^2) \quad (2)$$

$$\Rightarrow f(x_k|\theta) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_k - \theta)^2}{2\sigma^2}\right) \quad (3)$$

$$\Rightarrow L(\theta) = (2\pi\sigma^2)^{-\frac{T}{2}} \prod_{k=1}^T \exp\left(-\frac{(x_k - \theta)^2}{2\sigma^2}\right) \quad (4)$$

$$= (2\pi\sigma^2)^{-\frac{T}{2}} \exp\left(-\frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2\right) \quad (5)$$

$$\Rightarrow l(\theta) = -\frac{T}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2 \quad (6)$$

Now, suppose θ_{ML} solves the following, $\frac{dl}{d\theta} = 0$ (7)

$$= -\frac{1}{\sigma^2} \sum_{k=1}^T (x_k - \theta_{\text{ML}}) \quad (8)$$

$$\Rightarrow 0 = \sum_{k=1}^T (x_k - \theta_{\text{ML}}) \quad (9)$$

$$\Rightarrow \theta_{\text{ML}} = \frac{1}{T} \sum_{k=1}^T x_k = \bar{x} \quad (10)$$

Numerical Maximum Likelihood

```
import numpy as np
from matplotlib import pyplot as plt
from scipy.optimize import minimize_scalar
from scipy.stats import norm
```

```

# -- Plotting Tools --
def low_buff(y_values):
    span = y_values.max() - y_values.min()
    return y_values.min() - 0.1 * span

# -- True Distribution --
THETA = 1
SIGMA = 1

# -- Samples
T = 1000

r = np.random.normal(scale=SIGMA, size=(T,))
sample = THETA + r
print(sample.mean())

## 1.0138583286773388


def neg_log_llh(theta, sample, T, sigma):
    return - T/2 * np.log(2 * np.pi * sigma**2) + 1 / (2 * sigma) * ((sample - theta)**2).sum()

def nll_vec(theta_plot, sample, T, sigma):
    return np.array([neg_log_llh(theta, sample, T, sigma) for theta in theta_plot])

theta_plot = np.linspace(0, 3, 1000)
likelihood_plot = nll_vec(theta_plot, sample, T, SIGMA)
estimator_result = minimize_scalar(neg_log_llh, args=(sample, T, SIGMA))

fig, ax = plt.subplots()
ax.plot(theta_plot, likelihood_plot,
        label='Likelihood')
ax.vlines(x=THETA, ymin=low_buff(likelihood_plot), ymax=likelihood_plot.max(),
        label='True Parameter')
ax.plot(sample.mean(), neg_log_llh(sample.mean(), sample, T, SIGMA), 'o',
        label='ML Estimator\n' + r'$\theta_{\mathrm{ML}}$ = ' + str(round(sample.mean(), 3)) + '$')
# ax.plot(estimator_result.x, estimator_result.fun, 'o',
#         label=f'Numerical Minimization: {round(estimator_result.x, 3)}, '
#         f' {round(estimator_result.fun, 3)}')
ax.set_xlabel(r'$\theta$')
ax.set_ylabel('Likelihood ' + r'$L(\theta)$')
ax.set_title('Likelihood Function')
ax.legend()
plt.show()

```



Analytical Maximum A Posteriori

This is Exercise 11.2 from (Särkkä and Solin 2019). Since this is not a nonlinear SDE, we can use the following to determine the posterior probability distribution:

$$p(\theta|x(t_1), \dots, x(t_T)) \sim p(\theta) \exp(-l(\theta)) = \exp(-l_p(\theta)) \quad (11)$$

$$= \left[\frac{1}{\sqrt{2\pi\lambda^2}} \exp\left(-\frac{\theta^2}{2\lambda^2}\right) \right] \left[\exp\left(\frac{T}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2\right) \right] \quad (12)$$

$$= (2\pi\lambda^2)^{\frac{1}{2}} \exp\left(-\frac{\theta^2}{2\lambda^2} + \frac{T}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2\right) \quad (13)$$

and to find the maximum a posteriori estimate, $\theta_{\text{MAP}} = \arg \min_{\theta} p(\theta|x(t_1), \dots, x(t_T))$, we minimize the

negative logarithm of the posterior distribution, which we'll call A .

$$A = -\log \left((2\pi\lambda^2)^{\frac{1}{2}} \exp \left(-\frac{\theta^2}{2\lambda^2} + \frac{T}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2 \right) \right) \quad (14)$$

$$= -\frac{1}{2} \log(2\pi\lambda^2) + \frac{\theta^2}{2\lambda^2} - \frac{T}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{k=1}^T (x_k - \theta)^2 \quad (15)$$

$$= -\frac{1}{2} \log(2\pi\lambda^2) + \frac{\theta^2}{2\lambda^2} - \frac{T}{2} \log(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{k=1}^T x_k^2 - 2\theta \frac{1}{2\sigma^2} \sum_{k=1}^T x_k + \frac{1}{2\sigma^2} T\theta^2 \quad (16)$$

$$= \left[\frac{1}{2\lambda^2} + \frac{T}{2\sigma^2} \right] \theta^2 - \left[\frac{1}{\sigma^2} \sum_{k=1}^T x_k \right] \theta + \left[\frac{1}{2\sigma^2} \sum_{k=1}^T x_k^2 - \frac{1}{2} \log(2\pi\lambda^2) \right] \quad (17)$$

Since A is quadratic in θ , we can use the familiar formula for the minimum.

$$\theta_{\text{MAP}} = \frac{\frac{1}{\sigma^2} \sum_{k=1}^T x_k}{\frac{1}{\lambda^2} + \frac{T}{\sigma^2}} \quad (18)$$

$$= \frac{\frac{1}{\sigma^2} \sum_{k=1}^T x_k}{\frac{T\lambda^2 + \sigma^2}{\lambda^2\sigma^2}} \quad (19)$$

$$= \frac{\lambda^2}{T\lambda^2 + \sigma^2} \sum_{k=1}^T x_k \quad (20)$$

$$= \frac{1}{T + \frac{\sigma^2}{\lambda^2}} \sum_{k=1}^T x_k \quad (21)$$

Numerical MAP Estimator

```
# -- Prior --
LAMBDA = 2

def prior(theta, lamb):
    return 1 / np.sqrt(2 * np.pi * lamb**2) * np.exp(- theta**2 / (2 * lamb**2))

def posterior(theta, lamb, sigma, sample):
    """ f1 * exp(f2 + f3 + f4) """
    f1 = np.sqrt(2 * np.pi * lamb**2)
    f2 = - theta**2 / (2 * lamb**2)
    f3 = sample.size / 2 * np.log(2 * np.pi * sigma**2)
    f4 = - 1 / (2 * sigma**2) * np.sum((sample - theta)**2)
    return f1 * np.exp(f2 + f3 + f4)

def map_estimate(T, sigma, lamb, sample):
```

```

""" Uses analytical formula found by minimizing the
un-normalized negative log of the posterior """
return 1 / (T + sigma**2/lamb**2) * sample.sum()

def this_neg_posterior(theta):
    """ dependent on current samples
    convenience function for optimizer of the posterior
    without mixing arrays (i.e. sample size vs attempted theta) """
    return -1 * posterior(theta, LAMBDA, SIGMA, sample)

# Find Estimators
num_map_estimate = minimize_scalar(this_neg_posterior)
theta_map = map_estimate(T, SIGMA, LAMBDA, sample)

# plotting
theta_plot = np.linspace(-5, 5, 5000)
prior_plot = prior(theta_plot, LAMBDA)
posterior_plot = np.array([posterior(th, LAMBDA, SIGMA, sample) for th in theta_plot])

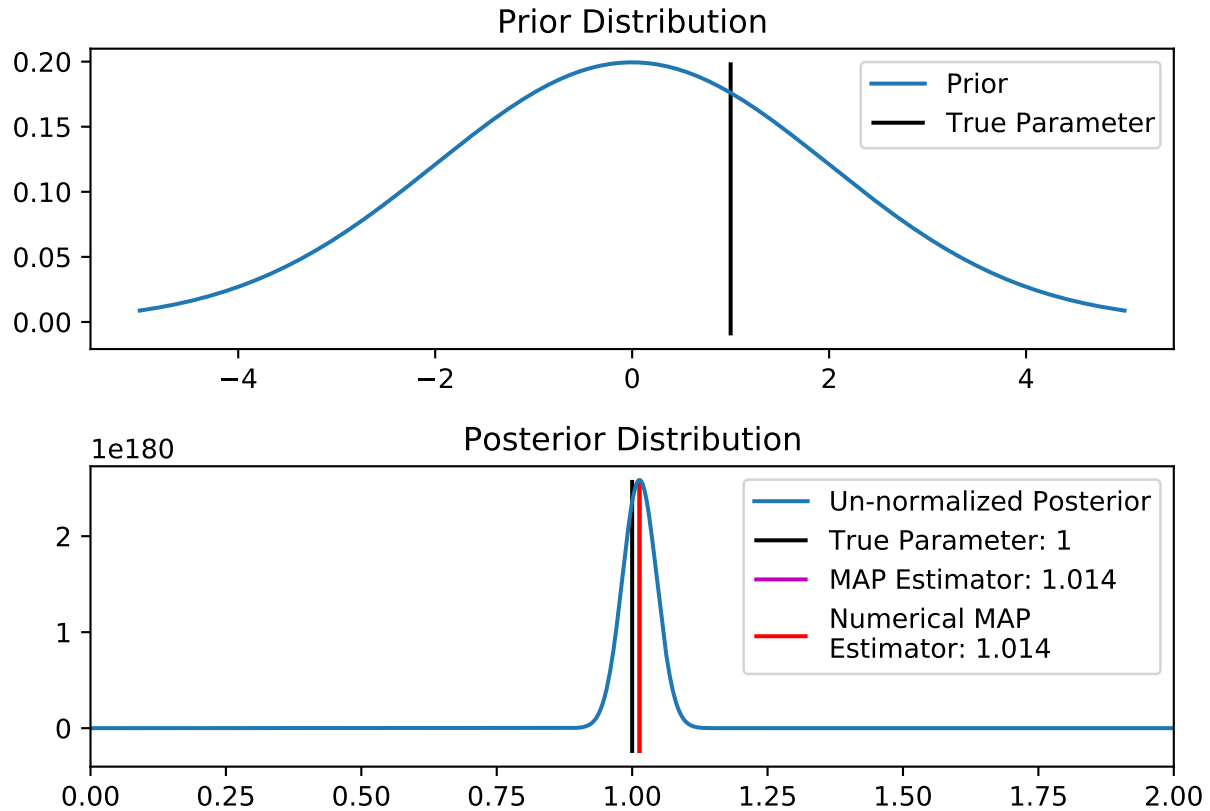
fig2, axes2 = plt.subplots(nrows=2)
# prior plot
axes2[0].plot(theta_plot, prior_plot, label='Prior')
axes2[0].vlines(x=THETA, ymin=low_buff(prior_plot), ymax=prior_plot.max(),
               label='True Parameter')
axes2[0].set_title('Prior Distribution')
axes2[0].legend()

# posterior plot
# Doesn't make sense to show prior since posterior is un-normalized
# axes2[1].plot(theta_plot, prior_plot, label='Prior', ls='dashed', alpha=0.8)
axes2[1].plot(theta_plot, posterior_plot, label='Un-normalized Posterior')
axes2[1].vlines(x=THETA, ymin=low_buff(posterior_plot), ymax=posterior_plot.max(),
               label=f'True Parameter: {THETA}')
axes2[1].vlines(x=theta_map, ymin=low_buff(posterior_plot), ymax=posterior_plot.max(),
               colors='m', label=f'MAP Estimator: {round(theta_map, 3)}')
axes2[1].vlines(x=num_map_estimate.x, ymin=low_buff(posterior_plot), ymax=posterior_plot.max(),
               colors='r', label=f'Numerical MAP\nEstimator: {round(num_map_estimate.x, 3)}')
axes2[1].set_xlim(left=0, right=2)

## (0.0, 2.0)

axes2[1].set_title('Posterior Distribution')
axes2[1].legend()
fig2.tight_layout()
plt.show()

```



Metropolis-Hastings Algorithm

This is Exercise 11.3 from (Särkkä and Solin 2019). This is strictly a coding exercise, using the Metropolis-Hastings algorithm to approximate the posterior distribution.

```
# Functions for M-H Algorithm
def sample_proposal(cur_theta, gamma):
    return np.random.normal(loc=cur_theta, scale=gamma)

def proposal_probability(prop_theta, loc, gamma):
    return norm(loc=loc, scale=gamma).pdf(prop_theta)

def l_p(theta, lamb, sigma, sample):
    """ Notes on derivation
        l(th) is the negative log-likelihood
        p(th) is the prior distribution

        l_p = l(th) - log(p(th))
            = - log( p(th) * exp(-l(th)) )

        written in the form f1*theta**2 + f2 * theta + f3 + f4
    """
```

```

T = sample.size
f1 = 1 / (2 * lamb**2) + T / (2 * sigma**2)
f2 = - 1 / sigma**2 * (sample.sum())
f3 = 1 / (2 * sigma**2) * np.sum((sample**2))
f4 = - 1/2 * np.log(2 * np.pi * lamb**2)
return f1 * theta**2 + f2 * theta + f3 + f4

def acceptance_probability(cur_theta, prop_theta, lamb, sigma, sample, gamma):
    ratio_1 = np.exp(l_p(cur_theta, lamb, sigma, sample)
                    - l_p(prop_theta, lamb, sigma, sample))
    ratio_2 = (proposal_probability(cur_theta, prop_theta, gamma)
              / proposal_probability(prop_theta, cur_theta, gamma))
    return min(1, ratio_1 * ratio_2)

def metropolis_hastings(theta_0, gamma, lamb, sigma, sample, num_steps):
    mcmc_sampling = [theta_0]
    cur_theta = theta_0
    accept_count = 0
    for _ in range(num_steps):
        prop_theta = sample_proposal(cur_theta, gamma)
        alpha = acceptance_probability(
            cur_theta, prop_theta, lamb, sigma, sample, gamma
        )
        u = np.random.uniform()
        if u < alpha: # accept
            mcmc_sampling.append(prop_theta)
            cur_theta = prop_theta
            accept_count += 1
        else:
            mcmc_sampling.append(cur_theta)
    acceptance_rate = accept_count / num_steps
    print(f'Acceptance Probability: {acceptance_rate}')
    return mcmc_sampling, acceptance_rate

# Metropolis-Hastings Values
MCMC_TRIALS = 10000
THETA_0 = 0
GAMMA = .15

mh_sampling, acceptance_rate = metropolis_hastings(
    THETA_0, GAMMA, LAMBDA, SIGMA, sample, MCMC_TRIALS
)

# construct analytical posterior normalized:

## Acceptance Probability: 0.2531

mh_theta_plot = np.linspace(
    np.min(mh_sampling),
    np.max(mh_sampling),

```

```

    1000,
)
mh_analytical_post = np.array([posterior(th, LAMBDA, SIGMA, sample) for th in mh_theta_plot])
tmp_hist = np.histogram(mh_sampling, 'auto')
peak = tmp_hist[0].max()
del tmp_hist
mh_analytical_post *= peak / mh_analytical_post.max()

# Plot mh_sampling over time
mh_fig, mh_axes = plt.subplots(nrows=2)
mh_axes[0].plot(mh_sampling)
mh_axes[0].set_xlabel('Iteration')
mh_axes[0].set_ylabel(r'$\theta$')
mh_axes[0].set_title(f'Sampling History\nAcceptance Rate: {round(acceptance_rate, 3)}')

# Plot distribution
_ = mh_axes[1].hist(mh_sampling, bins='auto', label='MCMC')
mh_axes[1].plot(mh_theta_plot, mh_analytical_post, label='Analytical', ls='dashed')

## [<matplotlib.lines.Line2D object at 0x120461880>]

mh_axes[1].vlines(x=THETA, ymin=0, ymax=mh_analytical_post.max(),
                 label=r'True Parameter: $\theta$ = ' + f'{THETA}')

## <matplotlib.collections.LineCollection object at 0x11dfe1640>

mh_axes[1].set_xlabel(r'$\theta$')

## Text(0.5, 0, '$\theta$')

mh_axes[1].set_title(r'Posterior Distribution (normalized)')

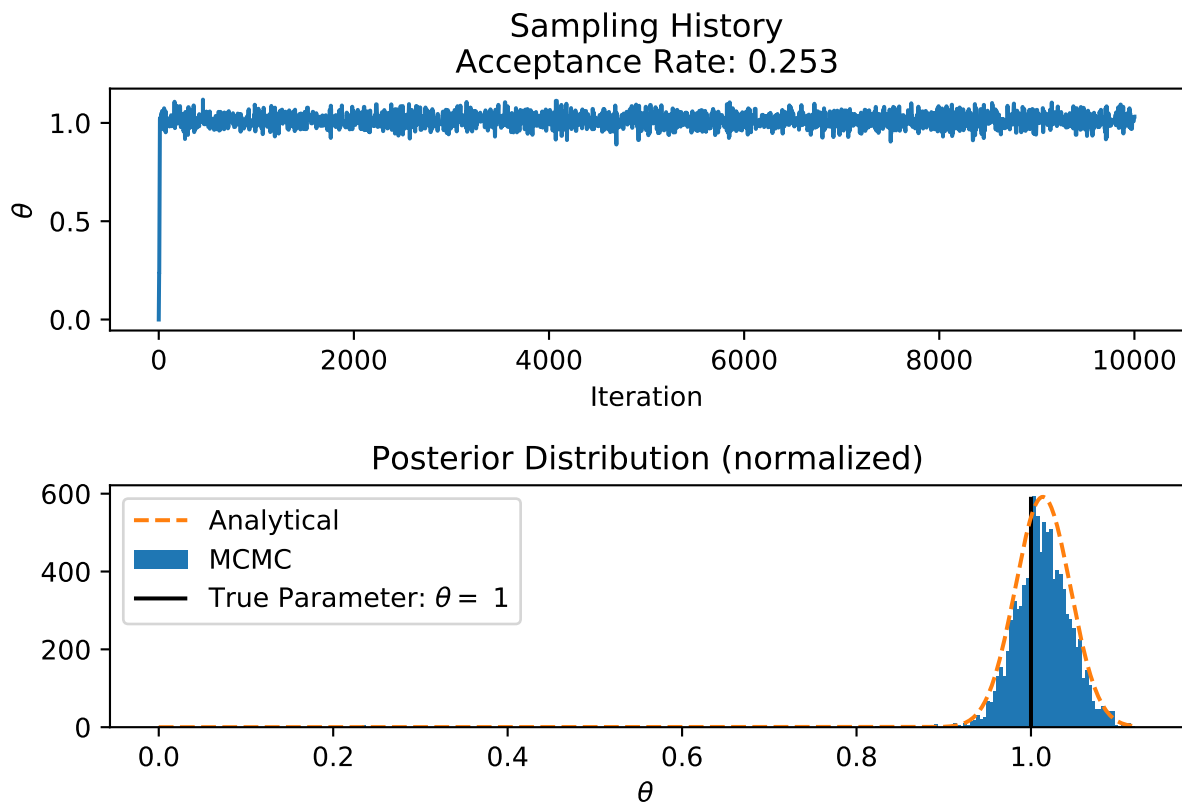
## Text(0.5, 1.0, 'Posterior Distribution (normalized)')

mh_axes[1].legend()

## <matplotlib.legend.Legend object at 0x11dfe1f40>

mh_fig.tight_layout()
plt.show()

```

References

Särkkä, Simo, and Arno Solin. 2019. *Applied Stochastic Differential Equations*. Vol. 10. Cambridge University Press.